Dynamic Load Balancing for Heterogeneous Systems

Kenneth Lee, Kyle Morgan

Abstract—**Leveraging maximum performance from specialized hardware is one of the major problems facing computer scientists today. General-purpose computing on graphics processing units (GPGPU) is a technique that is becoming increasingly popular for handling computations typically performed by the CPU. However, many GPGPU algorithms place the entire workload on the Graphics processing unit (GPU) leaving the CPU idle for the duration of the computation. This paper investigates the effect various load-balancing schemes across the CPU and GPU have on overall application performance. Additionally, the discrete and fused GPU architectures, and their impact on performance, will be investigated.**

I. INTRODUCTION

LEVERAGING the maximum performance from specialized hardware is one of the major problems facing computer scientists today. The unique architecture of the GPU gives it a place as a massively parallel hardware accelerator. Programming frameworks such as OpenCL allow programmers to utilize the stream processors of the GPU for non-graphics data. However, because of the difficulty in programming for this hardware, most of the effort is spent on trying to achieve performance from the GPU, while the CPU is left unused. By also making use of CPU to process a chunk of the data, we can achieve higher utilization of resources and therefore increase overall program performance.

GPU+CPU co-processing hardly a new idea. Jimenez et al. [8] present a dynamic library which can be used in a system to dynamically schedule tasks for GPU computation across various processes. The method described only looks at optimizing usage of the GPU across the whole system, rather than optimizing a single applications performance. The StarPU system, on the other hand, presents a method of co-processing and load balancing on a per application basis [2]. They show that by using an appropriate model for the load balancing, they are able to achieve improvements in speed for applications which call for code acceleration multiple times. We intend to extend this work to improve the speed of each kernel execution by splitting the work of that kernel over the capable resources, similar to the schemes used by OpenMP [5]. Both of these papers address data transfer times as a major cost of GPU computing. Using the new AMD Fusion architecture, which fuses the CPU and GPU onto a single die, we may be able to eliminate or greatly reduce this cost, which will impact the overhead of GPU+CPU co-processing.

We expect that by using GPU+CPU processing we will improve the performance of the application, when compared to CPU-only or GPU-only. That is, consider a workload of size n, a GPU that requires t_g time to perform one block of computation, and a CPU that requires t_c time to perform one block of computation. We expect that the GPU would require time nt_g to perform a given computation, while the CPU would require nt_c time to perform the same computation. Optimally, we would expect the load-balancing scheme to require the following time to perform the computation.

$$\max\left\{(n-m)t_g, mt_c\right\} + t_d + t_o$$

Where m is the amount of the workload given to the CPU ($m \leq n$), t_d is the data transfer cost, and t_o is any overhead incurrent from synchronization costs of the load balancing scheme or otherwise. The above will be faster than just the GPU iff $t_g - \max\left\{(n-m)t_g, mt_c\right\} < t_d + t_o$.

Two different load balancing schemes were implemented for this paper. The first was a static load balancer, which partitions the workload in two, and transfers one chunk to the CPU, and the other to the GPU. The CPU and GPU work entirely independently, and their results are merged once each finish. The static load balancer should have no synchronization overhead, and transfers all of the data prior to the computation. The main caveat with a static load balancer is that the ratio of the workload given to the CPU is chosen by the load balancer. An optimal ratio would be one such that the time to perform the computation is equal for both the CPU and GPU. Unfortunately, this optimal ratio will be application-dependent, meaning that a fixed ratio cannot be optimal for every application. In order for a static load balancer to be optimal for a given application, a programmer must tune the amount of workload given to the CPU that works best for that application.

The second scheme is a dynamic load balancer, which partitions the data into many smaller chunks. One chunk is sent to the CPU and GPU each for computation. When one finished, it requests another chunk from the load balancer, and this process repeats until the entirely workload has been sent to either the CPU or GPU. The dynamic load balancer will incur some amount of synchronization overhead, because shared variables that track which chunks of the workload are left must be protected by a lock. Unlike the static load balancer, the dynamic load balancer transfers data throughout the computation. The dynamic load balancer is designed to mitigate the primary problem with the static load balancer – that the optimal ratio of computation performed on the CPU is application-dependent. A chunk size must be chosen such that the processing units are not requesting chunks so frequently that contention for shared variables rise and synchronization overhead increases, yet small enough to not

allow one processing unit to receive a disproportionate amount of work.

II. ARCHITECTURE OVERVIEW

In this section we discuss the different architectures of the experimental machines under test. We will discuss CPU, GPU, and APU architectures in terms of major differences that affect the performance of the device.

A. CPU Architecture

The CPU architecture is designed as a low latency architecture. To facilitate the needs for low latency application performance, a series of hierarchical caches have been designed to hide memory access costs, the most common cause of latency in single-threaded applications on modern architectures. Each core of the CPU typically only runs one or two hardware-enabled threads to reduce cache pollution, thus avoiding latency. An simplistic example CPU architecture is given by Figure 1.

The use of the Single Instruction Multiple Data (SIMD) paradigm enables increased performance on these devices by increasing data throughput for data parallel applications. However, optimal performance of these types of instructions often involves complex compilation techniques and hand-tuning of the assembly.

B. GPU Architecture

The GPU almost acts a foil to the CPU architecture. Instead of optimizing the architecture of the device to increase performance of a single thread, the GPU architecture is specifically designed for large throughput performance. Large numbers of threads can be easily created an run on a GPU. Large register files enable tens of thousands of threads to be run on today's GPU architecture with very little penalty for context switching.

Groups of processing units, Streaming Multiprocessors (SMs), share a single instruction dispatch unit and operate in lock step. The threads themselves run on each SM in a group, called a warp, typically of 32 threads. Because of these architectural features divergent branching in a warp (caused by if-else clauses) greatly degrade application performance.

The memory architecture for the GPU is very different than that found on the CPU. Instead of a large hierarchy of caches to hide memory latency for single thread performance the latency to memory is instead hidden through aggressive context switching. Newer GPU architectures may contain privative caching schemes. Each SM on the GPU also contains a local memory memory space which can accessed at much lower latency . This memory space is typically used as a user-managed cache. The memory hierarchy for the GPU architecture is given by Figure 2.

C. APU Architecture

The APU is a relatively new, novel architecture that attempts to improve on the the discrete CPU+GPU system. The CPU die is split between the typical CPU processor as well as allocating space for a SIMD engine, the basis for a GPU


Fig. 1. CPU Architecture

Fig. 2. GPU Architecture

processor. The APU has a shared memory between the CPU and SIMD engine, in which accesses are facilitated through a high-speed memory controller. The AMD Fusion technology is an example of this type of architecture [1].

One of the main benefits of the APU is the elimination of the PCI-e bus, which allows for much faster transfer between CPU and GPU. Daga et al. show that the AMD Fusion architecture allows for much faster data transfers between CPU and GPU, allowing for an overall improvement in runtime, despite reduced compute power on the GPU [4].

III. LOAD BALANCING BENCHMARKS

For this paper, three different benchmarks were used to assess the effectiveness of the load balancer. These benchmarks were chosen to represent tasks that would actually be performed on a GPU, and have seen use in already existing benchmark suites [6]. The applications that we have chosen for this project are VectorAdd, VectorAdd+, and Reduce.

The VectorAdd application is a simple vector addition operation that computes the sum of two input vectors, and stores the result into an output vector. This represents the vector operation $\vec{C} = \vec{A} + \vec{B}$. This kernel is naive to parallelize, as it exhibits embarrassingly parallel structure. Each thread being run maps to a single index in the array and computes the output value for that index. This algorithm is also nicely designed for load-balancing because of the lack of data dependencies, reducing the amount of communication to the GPU.

The next benchmark, referred henceforth as the "Vector Add+" benchmark, is similar to the VectorAdd benchmark, but it computes 1000 additions before storing the result in the output vector. The motivation behind this benchmark is to reduce memory access overhead with more computation. This

benchmark was compiled without optimizations so that the additions would not be optmized into a single multiplication. The data dependencies are exactly the same as for the VectorAdd application, so the results for data transfer speeds should be equal for both applications, and the difference between the two should only be the kernel execution times.

The Reduce application is an application that performs a reduction operation on a large array of data. In this case, the sum of a variably-sized array is computed. The Reduce benchmark is unique in that different code is being run on the CPU and the GPU. The reason for this is that the GPU cannot efficiently compute a reduction the way a CPU typically would – due to the relatively slower cores and much greater memory latency. A optimized reduction on the GPU is a multi-stage process, in which the output of one step is used as the input of the next. Each SM computes the reduction of a smaller set of elements from the larger array during each iteration. This requires the use of numerous local atomic operations, as well as local memory to facilitate optimized memory access between threads on an SM. This process continues until the reduction operation is only done by a single SM. A natural consequence of this algorithm is that not all of the SMs will be utilized during the computation. In addition, inside an iteration all of the streaming processors will only be used for the first stage of reduction, greatly reducing utilization on the device.

All numbers reported in the results section are given as a mean of ten runs for that particular benchmark, load-balancing scheme, and data size.

IV. THE OPENCL FRAMEWORK

For this project, we will use the OpenCL Framework for our kernel lauching. OpenCL is an application framework for heterogeneous parallel computing. It includes an API for creating and transferring memory to accleration devices, as well as dynamic compilation of application kernels that are run on the acceleators.Using OpenCL allows us to write the host code only once, while allowing us to run the core application kernel on the CPU, GPU, or both.

In OpenCL each thread is logically mapped to a work-item. These work-items are grouped into workgroups. These workgroups map to a SM on a discrete GPU. For instance, these workgroups share access to a SMs local memory and can perform a barrier with all the other work-items in the workgroup. In the case of a CPU acclerator, the workgroups still exist, but optimizations such as manual caching via local memory will not be effective.

Devices are grouped into contexts, which contains memory buffers and so on. A queuing system is used by OpenCL to facilitate kernel execution. Kernels are submitted to the queue along with the size of each workgroup, as well as number of workgroups to execute. This queue is also responsible for reading and writing memory buffers. The queue is in-order, meaning that the queue will wait until previously enqueued items have finished before starting work on more recently queued items.

For our purposes, we decided to use a seperate context for each device, as well as a seperate command queue for each context. We decided to use a seperate context for each device, as it would allow for finer-grained control of data transfers and residency. Using a seperate command queue for each device will facilitate multiple devices to be used simultaneously, as each command queue operates on its own seperate thread.

V. LOAD BALANCING SCHEMES

A naive implementation of CPU+GPU co-processing will in many cases lead to an imbalance of work for either the CPU or GPU. To deal with this problem we will investigate the use of 2 different load balancing schemes in addition to the performance of a CPU-only or a GPU-only solution. These schemes are illustrated in Figure 3.

The first load balancing scheme that was implemented was a static division of work between the CPU and the GPU at a specified ratio. This ratio was user-defined allowing for best performance after runtime benchmarking (i.e. the precentage of work done on the CPU and GPU). The cases of CPU-only and GPU-only scheduling are just special cases of this scheduler, of the ratios 0 and 1 respectively.

The static load-balancing scheme benefits from the low overhead of only have to launch the kernels on each device only once, as opposed to more complex load-balancing schemes. This increases the amount of work to be done at one time, as workgroups that retire are immediately replaced by newer workgroups from the larger queue. The main cause of overhead for this method of load-balancing is that of a varying workload. If the workload for the kernel is imbalanced between workgroups, it is impossible to statically determine the proper amount of work to put on each kernel. Improper allocation of work to each device will not allow both devices to be fully utilized, and application performance will not be optimal.

To address the issue of a dynamic work load between workgroups for kernels, we also used a dynamic load-balancing scheme. In this scheme the workgroups were grouped into larger chunks. Each device is initially allocated one chunk of work to work on, and as each chunk is retired by the device a new chunk is allocated. Each device is given a seperate hardware thread, and the thread will wait until the current chunk is finished, then it will synchronize through a shared variable to acquire the next chunk of work for the device. This process will continue until the work queue is out of work.

The dynamic load-balancing scheme will still not guarantee optimal performance. Within a work chunk, a single long-running workgroup can hold up the entire device, not allowing other kernels to execute on open SMs. This creates and issue where the size of the work chunks should be considered. One one hand, larger work chunks can provide for better performance for each device, as it will allow more workgroups to hide the cost of a long-running workgroup. However, large chunk sizes can reduce the granularity of load-balancing between the devices. Alternatively, smaller chunk sizes alleviates the problem of one device receiving a disproportionate amount of the work, but can increase scheduling overhead by increasing contention for the shared variable which tracks the remaining workload.

Fig. 3. Data Partitioning Schemes for Load Balancing.

Fig. 5. Vector Add Benchmark on a Discrete GPU Platform

Fig. 4. Reduction Benchmark on a Discrete GPU Platform

Fig. 6. Vector Add+ Benchmark on a Discrete GPU Platform

VI. RESULTS AND ANALYSIS

We tested the different benchmark applications using different hardware environments. For the discrete GPU tests, we tested our application on an Intel Core Duo CPU with a AMD HD5450 GPU attached via PCIe. For the Fused CPU+GPU environment we used an AMD Fusion A-3850. Linux was used for the discrete environment, while Windows was used on the Fused environment in order to take advantage of newer drivers that allowed for better hardware features on the new Fusion machine.

The three implemented benchmarks were run on a system with a discrete GPU, as well as on a fused GPU platform. In dynamic load balancing benchmarks, a chunk size of 80 KB was used. Dynamic results are only given for the discrete GPU testing environment, as the APU platform was running Windows and our code was written with POSIX compliance in mind. In static load balancing benchmarks, the ratio of data computed on the GPU will be specified.

First, the reduction benchmark was run on the discrete GPU platform with a variety of different data sizes and load balancing policies, as shown in Figure 4. Interestingly, the CPU outperformed the GPU on this particular benchmark. This is likely due to the GPU not utilizing all of its stream processors during a reduction operation, as discussed previously. It is also the case that the large amount of synchronization that must occur in each iteration of the GPU reduction, and the number of kernel launches required for each iteration that is not present in the CPU reduction. In addition, because the CPU had better performance on the kernel alone, it saw an even greater improvement when the data transfer costs were considered, with the CPU-only case requiring no data movement at all. Unfortunately, none of the static load balancing schemes were able to outperform the CPU-only scheme; all the static times were somewhere in between the CPU-only and GPU-

only times. Additionally, the dynamic load balancing scheme performed significantly worse than all the other schemes, and more than 100% worse than just the CPU. The cause for this degradation in performance was likely due to the overhead of scheduling threads and communciation between those threads, as well as moving data to the GPU in small chunks when the GPU showed very poor performance compared to the CPU.
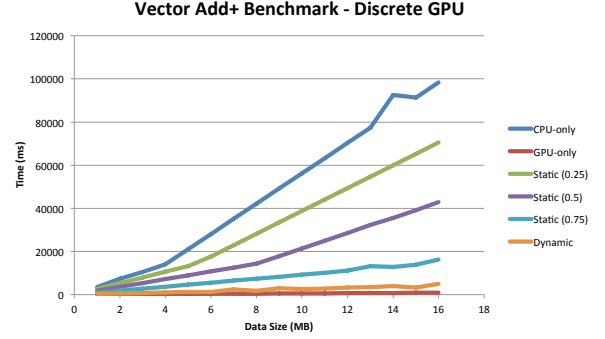
Now, consider Figure 5, where the VectorAdd benchmark was run on the discrete GPU platform. In this case, we see that the GPU has nearly a 4x speedup over just the CPU. The data parallelism and lack of any synchronization requirements makes this speedup over the CPU possible. However, similar to the reduction benchmark, the static load balancing schemes never outperform just the GPU-only scheme, with all of them lying somewhere between the CPU-only and GPU-only runs. Also similar to the reduction benchmark, the dynamic load balancing scheme performs significantly worse than all the other schemes.

The results for the VectorAdd+ benchmark on the discrete GPU platform are shown in Figure 6. There are a few notable trends in this figure that are not present in the previous two. The first is is how drastically the GPU outperforms the CPU – with over a 100x speedup over just the CPU with 16 MB of data. Because of the very large amount of computation that needs to be performed, compared to the low amount of data accesses required in the kernel (3 data accesses for 1000 additions), the GPU is able to outperform the CPU because of its very large throughput for computationally intensive kernels. The second thing to notice for this benchmark specifically is that the dynamic load balancing drastically outperforms the static load balancing schemes. This is almost certainly

Fig. 7.   Varying Data Ratio with a Discrete GPU

Fig. 9.   Vector Add Benchmark on a Fused CPU+GPU Platform


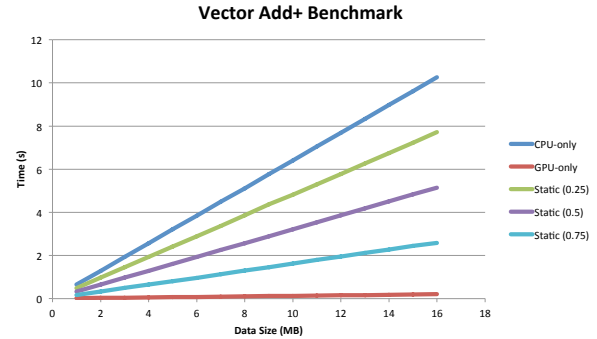Fig. 8.   Reduction Benchmark on a Fused CPU+GPU Platform

Fig. 10.   Vector Add+ Benchmark on a Fused CPU+GPU Platform

because the GPU so drastically outperforms the CPU. With the dynamic load balancing scheme, only small chunks of data are given to the CPU at a time, rather than a static percentage, allowing for the GPU to take advantage of much more of its speedup over the CPU. Increasing the chunk size for the GPU (taking two chunks instead of one at a time) would likely yield in a greater speedup for the dynamic load-balancing case.

For the above figures, the static load balancer was run with three different data ratios run on the GPU – 25%, 50%, and 75%. The general trend was that as more of the workload was placed on the GPU, the runtime improved, with an exception in the Reduce benchmark due to the superior CPU performance. As a result, we run the vector add benchmark again with a high granularity of data ratios. Figure 7 shows this benchmark on the discrete GPU platform with between 90% and 100% of the data on the GPU, with a 0.1% granularity. The total execution time, including data transfers, has an obvious downward trend as more of the data is executed on the GPU. That said, there are a couple of static ratios which yielded better performance than executing all of the data on the GPU. Most notably, when only 98.6% of the data was executed on the GPU. In this instance, the overall execution time was 325.27 seconds, compared to 333.31 seconds with only the GPU – a 2.4% speedup. However, these few instances of speedup may simply be due to variability in the system.

Next, consider Figures 8, 9, and 10. The figures demonstrate a variety of different load balancing options for the Reduction, VectorAdd, and VectorAdd+ benchmarks, across data sizes ranging between 1 MB and 16 MB on a Fused CPU+GPU platform. Each figure depicts the total execution time for the CPU-only and GPU-only computations, as well as static load

balancing with 25%, 50%, and 75% of the data on the GPU. In each of the figures, it's evident that the GPU performs significantly faster than the CPU for the given benchmarks, with the static partitioning schemes lying somewhere between the two. As you can see in the figures, while the overall performance numbers have changed between the discrete and fused architectures the major trends from the previous graphs still hold true.

The only deviation from the above is the reduction benchmark, in which the GPU only performs   14% faster than the CPU in the 16 MB case (0.104 seconds for the GPU, 0.091 seconds for the CPU), and some of the static load balancing schemes perform slightly better than just the GPU. For instance, in the 9 MB and 10 MB cases, both the 50% and 75% static schemes performed better than just the GPU (0.5% and 4.4% respective speedups for 9 MB, 1.5% and 3.6% respective speedups for 10 MB). What is also interesting to note is that on the fused CPU+GPU architecture, the GPU outperformed the CPU on the reduction benchmark. However, the opposite was true of the reduction benchmark on the discrete GPU platform. The difference in this performance was likely due to the increased transfer speeds between the CPU and GPU on the fused platform. It is also the case that the SIMD engine on the Fusion APU is slightly more powerful than the discrete GPU that we have used for testing, allowing kernel execution times on the APU to be lower than that of the discrete GPU.

Finally, consider Figure 12, where the overall runtime for the static load balancer is compared against both the CPU-only and GPU-only runtimes for 0.1%, 1%, 10%, 25%, 50%, 75%, 99%, and 99.9% static data ratios for the 16 MB vector

Fig. 11.   Reduction Benchmark with Varied Execution Ratio

Fig. 13.   Vector Add+ Benchmark with Varied Execution Ratio

Fig. 12.   Vector Add Benchmark with Varied Execution Ratio

Fig. 14.   Discrepancies between expected runtimes and actual runtimes.

add benchmark. Figures 11 and 13 show similar results for the reduction and vector add+ benchmarks. In these benchmarks, the downward trend as more data is executed on the GPU becomes more evident. What's also interesting is that the static load balancer actually performs worse than the CPU-only case when only a small amount of the data ($\leq 1\%$) is executed on the CPU in the reduction and vector add benchmarks.

Based on the times shown in the previous figures, it is clear that the static load balancing is ineffective at improving the overall efficiency of any of the benchmarks tested. There were a few scenarios in which a specific static execution ratio achieved meager speedup ($<5\%$) over just the GPU, but these situations were scarce, and did not achieve the speedup we desired. Additionally, this small amount of speedup would require a significant amount of tine-tuning by a programmer.

Consider the 16 MB vector add benchmark, where the CPU-only case took 0.07465 seconds to run, and the GPU-only case took 0.02873 second to run. Ideally, 72.2% of the workload would be placed on the GPU, and 27.8% on the CPU. Assuming no other overhead, this would result in an overall runtime of 0.02075 seconds – a 27.8% speedup over the GPU. Using similar reasoning, a 43.6% theoretical speedup is attainable for the reduction benchmark. A measly 2% would be expected from the vector add+ benchmark, because of the vast disparity in the CPU-only and GPU-only runtimes.

The above results have lead us to believe that the CPU and GPU kernels are not running truly in parallel as we initially believed. As such, we used the OpenCL profiling capabilities to report how long each of the kernels took using the static load balancer with 95% of the data on the CPU. Figure 14 shows the discrepancy between the actual total runtime of the vector

add benchmark and the expected runtime of sequential calls to the kernels (that is, the sum of the CPU kernel runtime and the GPU kernel runtime). If the kernels were truly running in parallel, we would expect to see a negative discrepancy between the actual runtime and expected runtime. It was impossible for us to simply time the start and end times of the different kernels, because while that information was available for each kernel launch, it was kept as device-speciifc times. We could not reconcile the two different timing schemes, because of the different clock frequencies and the lack of a mapping from time from clock start for either device.

Many of the trials produced a small positive discrepancy. There were only two trials which produced a negative discrepancy; however, these trials were unusual in other ways as well. The overall runtime of the CPU kernel was rough 5.1 ms in all of the trials except these two, which were 36.6 ms and 31.6 ms, respectively. We believe that OpenCL is beginning the CPU kernel, scheduling the GPU kernel before the CPU kernel has completed, running the GPU kernel to completion, then scheduling the CPU kernel again and letting it run to completion. Thus, the profiler's record of the runtime of the CPU kernel would include the runtime of the GPU kernel.

With regard to the discrete GPU platform and the fused CPU+GPU platform, it seems that the fused platform has some advantages. For instance, consider Figures 4 and 8. Although we will not compare overall runtimes of the two benchmarks, we can see that the GPU of the fused platform performed better than the GPU of the discrete platform relative to their CPUs. That is, the CPU outperformed the GPU on the reduction benchmark on the discrete GPU architecture, but the converse was true of the fused platform. This may be due to the fact that

the fused platform shares memory between the CPU and GPU, which lowers data transfer times between the two devices. On the discrete GPU system, the GPU typically experienced significant data transfer times for the GPU, and the effect of this would be minimized on the fused platform.

VII. RELATED WORK

Ryoo et al. discuss optimization strategies for GPU versus CPU runtime systems [9]. However, their research is based on a GPU-only or CPU-only implementation instead of a heterogeneous approach. The work presented in [7] describes a runtime and compiled framework for improving memory access latency and reducing data transfer overhead for GPU kernels. Our work differs from this, because the kernel execution was still being performed solely on the GPU, whereas our work investigates kernel execution on multiple devices.

Computation via heterogeneous systems is a relatively new research area for high performance computing. The various frameworks and large number of different accelerator devices can make this quite difficult for application programmers to manage. The StarPU system [2] presents a framework to allow for programmers to write numerical codelets and have them automatically run on a heterogeneous system. Scheduling work for this system is one of the most difficult parts of the system, because of the various overheads for work-sharing schemes. We simplify this work, by looking at a two-device system, and focus on the performance characteristics of different scheduling scheme. Verner et al. use this approach to create a high-throughput low-latency heterogeneous encryption system [10].

Jimenez et al. [8] present a dynamic library which can be used in a system to dynamically schedule tasks for GPU computation across various processes. The method describes only looks at optimizing usage of the GPU across the whole system, instead of optimizing a single application's performance.

VIII. CONCLUSION

It is clear that our original hypotheses were not supported by our results. That is, we were not able to see consistent, significant speedup from any of our load balancing schemes. In many scenarios, the GPU simply outperformed the CPU, with the static load balancing schemes falling somewhere between the two, and the dynamic load balancer being far worse than anything else. There were only two notable exceptions to this trend. The first was the reduction benchmark on the discrete GPU platform, where the CPU outperformed the GPU due to the low stream processor utilization inherent in a GPU reduction implementation. However, the static load balancing schemes were still outperformed by the CPU in this case, so no overall speedup was gained. The second exception was the vector add+ benchmark on the discrete GPU platform, where the dynamic load balancer actually outperformed the tested static load balancing schemes. However, this was due to the vast disparity in the GPU and CPU runtimes for this particular benchmark, as giving any significant amount of data to the CPU would hurt performance. With the dynamic load balancer, the CPU will only compute a few blocks of data, rather than 25% or more with the static load balancer.

When analyzing more finely-grained static execution ratios, we found that, in the case where the GPU outperforms the CPU, there is always a downward trend in overall runtime as more data is executed on the GPU. There were a couple sparse instances where a static execution ratio performed better than using just the GPU, but we chalked these up to typical variability in runtimes due to their infrequency and minute performance improvements.

When analyzing the results, we considered the possibility that the CPU and GPU kernels were actually running in serial, which would explain the performance trends we were seeing in our static and dynamic load balancing schemes. An AMD profiler was used to find out the times each kernel took to compute, and the sum of these times was compared to the actual runtime of the benchmark. We found that the actual runtime of the benchmark was typically slightly more than the sum of the runtimes of each kernel, which seemed to indicate that the kernels were indeed running serially, with some overhead for data transfer times.

Based on the CPU-only and GPU-only runtimes of each benchmark, we were able to calculate an optimal static data partitioning, such that the computations on the CPU and GPU took equal time to complete. Based on the disparity of runtimes between the CPU-only and GPU-only results, the optimal runtimes ranged from a 2% to nearly a 45% speedup over just the GPU alone. Unfortunately, none of these theoretical results were supported in practice. It seems to be a limitation of OpenCL that it will not run kernels on different devices in parallel, which effectively renders our current load balancing schemes pointless.

Additionally, it seems that a fused CPU+GPU architecture may provide some performance benefits over just a discrete GPU architecture. Our results show that the GPU outperformed the CPU on the fused platform's reduction benchmark, but this was not the case for the discrete GPU platform. However, as the CPUs and GPUs differed significantly between the two testing environments, we cannot say definitively that one architecture will always be better than the other.

REFERENCES

[1] http://fusion.amd.com

[2] C. Augonnet, S. Thibault, R. Namyst, and P.A. Wacrenier, "StarPU: A Unified Platform for Task Scheduling on Heterogeneous Multicore Architectures" in *Proceedings of the 15th International Euro-Par Conference on Parallel Processing*, 2009, pp. 863-874.

[3] S. Che, M. Boyer, J. Meng, D. Tarjan, J.W. Sheaffer, S.H. Lee, and K. Skadron, "Rodinia: A benchmark suite for heterogeneous computing" in *Proceedings of the 2009 IEEE International Symposium on Workload Characterization (IISWC)*, 2009, pp. 44-54.

[4] M. Daga, A.M. Aji, and W. Feng, "On the Efficacy of a Fused CPU+GPU Processor (or APU) for Parallel Computing" in *Application Accelerators in High-Performance Computing (SAAHPC)*, 2011, pp. 141-149.

[5] L. Dagum and R. Menon, "OpenMP: An industry-standard API for shared-memory programming" in *IEEE Computational Science and Engineering* 1998, pp. 4655.

[6] A. Danalis, G. Marin, C. McCurdy, J.S. Meredith, P.C. Roth, K. Spafford, V. Tipparaju, and J.S. Vetter, "The Scalable Heterogeneous Computing (SHOC) benchmark suite" in *Proceedings of the 3rd Workshop on General-Purpose Computation on Graphics Processing Units*, 2010, pp. 63-74.

[7] J.A. Jablin, P. McCormick, and M. Herlihy, "Scout: High-Performance Heterogeneous Computing Made Simple" in *Parallel and Distributed Processing Workshops and Phd Forum (IPDPSW)*, 2011, pp. 2093-2096.

[8] V.J. Jimnez, L. Villanova, I. Gelado, M. Gil, G. Fursin, and N. Navarro, "Predictive Runtime Code Scheduling for Heterogeneous Architectures" in *Proceedings of the 4th International Conference on High Performance Embedded Architectures and Compilers*, 2009, pp. 19-33.

[9] S. Ryoo, C.I. Rodrigues, S.S. Baghsorkhi, S.S. Stone, D.B. Kirk, and W.W. Hwu, "Optimization Principles and Application Performance Evaluation of a Multithreaded GPU Using CUDA" in *Proceedings of the 13th ACM SIGPLAN Symposium on Principles and practice of parallel programming*, 2007, pp. 73-82.

[10] U. Verner, A. Schuster, and M. Silberstein, "Processing data streams with hard real-time constraints on heterogeneous systems" in *Proceedings of the international conference on Supercomputing*, 2011, pp. 120-130.

APPENDIX
TEAM MEMBER CONTRIBUTIONS

Kenneth Lee was responsible for load-balancing schemes in both design and implementation. He was also in charge of system integration and testing of the final applications on the appropriate hardware devices. Kenneth was also in charge of deep analysis and explanation of result data.

Kyle Morgan was in charge of kernel design and implementation. In addition, he was also in charge of result compilation and preliminary analysis. Kyle was also responsible for background readings in both GPGPU computation in general, as well as specifics in terms of both GPU architecture and the OpenCL framework.